

# Introduction to ML & DL

Shan-Hung Wu  
*shwu@cs.nthu.edu.tw*

Department of Computer Science,  
National Tsing Hua University, Taiwan

Machine Learning

# Outline

- 1 What's Machine Learning?
- 2 What's Deep Learning?
- 3 About this Course...
- 4 FAQ

# Outline

① What's Machine Learning?

② What's Deep Learning?

③ About this Course...

④ FAQ

# Prior vs. Posterior Knowledge

- To solve a problem, we need an algorithm
  - E.g., sorting

# Prior vs. Posteriori Knowledge

- To solve a problem, we need an algorithm
  - E.g., sorting
  - *A priori knowledge* is enough

# Prior vs. Posterior Knowledge

- To solve a problem, we need an algorithm
  - E.g., sorting
  - *A priori knowledge* is enough
- For some problem, however, we do not have the a priori knowledge
  - E.g., to tell if an email is spam or not
  - The correct answer varies in time and from person to person

# Prior vs. Posteriori Knowledge

- To solve a problem, we need an algorithm
  - E.g., sorting
  - *A priori knowledge* is enough
- For some problem, however, we do not have the a priori knowledge
  - E.g., to tell if an email is spam or not
  - The correct answer varies in time and from person to person
- Machine learning algorithms use the *a posteriori knowledge* to solve problems

# Prior vs. Posteriori Knowledge

- To solve a problem, we need an algorithm
  - E.g., sorting
  - *A priori knowledge* is enough
- For some problem, however, we do not have the a priori knowledge
  - E.g., to tell if an email is spam or not
  - The correct answer varies in time and from person to person
- Machine learning algorithms use the *a posteriori knowledge* to solve problems
  - Learnt from *examples* (as extra input)



# Example Data $\mathbb{X}$ as Extra Input

- Unsupervised:

$$\mathbb{X} = \{\mathbf{x}^{(i)}\}_{i=1}^N, \text{ where } \mathbf{x}^{(i)} \in \mathbb{R}^D$$

- E.g.,  $\mathbf{x}^{(i)}$  an email

# Example Data $\mathbb{X}$ as Extra Input

- Unsupervised:

$$\mathbb{X} = \{\mathbf{x}^{(i)}\}_{i=1}^N, \text{ where } \mathbf{x}^{(i)} \in \mathbb{R}^D$$

- E.g.,  $\mathbf{x}^{(i)}$  an email

- Supervised:

$$\mathbb{X} = \{(\mathbf{x}^{(i)}, \mathbf{y}^{(i)})\}_{i=1}^N, \text{ where } \mathbf{x}^{(i)} \in \mathbb{R}^D \text{ and } \mathbf{y}^{(i)} \in \mathbb{R}^K,$$


- E.g.,  $y^{(i)} \in \{0, 1\}$  a spam label

# General Types of Learning (1/2)

- *Supervised learning*: learn to predict the labels of future data points

$$X \in \mathbb{R}^{N \times D} :$$


$$\mathbf{y} \in \mathbb{R}^{N \times K} : [\mathbf{e}^{(6)}, \mathbf{e}^{(1)}, \mathbf{e}^{(9)}, \mathbf{e}^{(4)}, \mathbf{e}^{(2)}]$$

$$\mathbf{x}' \in \mathbb{R}^D :$$


$$\mathbf{y}' \in \mathbb{R}^K : ?$$

# General Types of Learning (1/2)

- **Supervised learning**: learn to predict the labels of future data points

$$X \in \mathbb{R}^{N \times D} :$$



$$y \in \mathbb{R}^{N \times K} : [e^{(6)}, e^{(1)}, e^{(9)}, e^{(4)}, e^{(2)}]$$

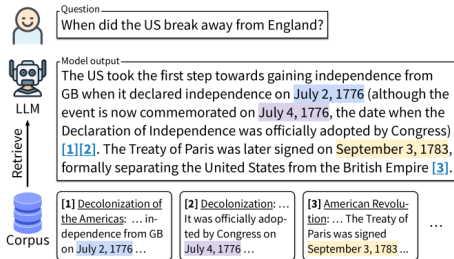
$$x' \in \mathbb{R}^D :$$



$$y' \in \mathbb{R}^K :$$

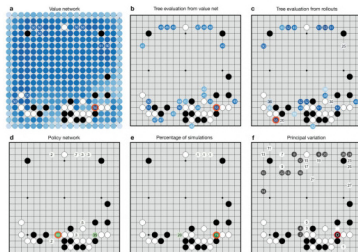
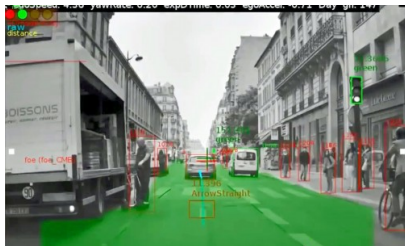
?

- **Unsupervised learning**: learn (latent) patterns in  $X$ , and optionally generate new  $x'$ s



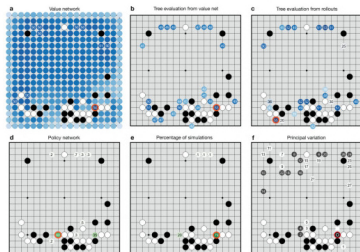
# General Types of Learning (2/2)

- **Reinforcement learning**: learn from “good”/“bad” feedback of actions (instead of correct labels) to maximize the goal



# General Types of Learning (2/2)

- **Reinforcement learning**: learn from “good”/“bad” feedback of actions (instead of correct labels) to maximize the goal



- AlphaGo [1] is a hybrid of reinforcement learning and supervised learning
  - Supervised learning from the game records
  - Then, reinforcement learning from self-play

# General Machine Learning Steps

- ① Data collection, preprocessing (e.g., integration, cleaning, etc.), and exploration

# General Machine Learning Steps

- ① Data collection, preprocessing (e.g., integration, cleaning, etc.), and exploration
  - ① Split a dataset into the training and testing datasets



# General Machine Learning Steps

- ① Data collection, preprocessing (e.g., integration, cleaning, etc.), and exploration
  - ① Split a dataset into the training and testing datasets
- ② Model development
  - ① Assume a **model**  $\{f(\cdot; \mathbf{w})\}$  that is a collection of candidate functions  $f$ 's (representing posteriori knowledge) we want to discover
    - ①  $f$  is assumed to be parametrized by  $\mathbf{w}$

# General Machine Learning Steps

- ① Data collection, preprocessing (e.g., integration, cleaning, etc.), and exploration
  - ① Split a dataset into the training and testing datasets
- ② Model development
  - ① Assume a **model**  $\{f(\cdot; \mathbf{w})\}$  that is a collection of candidate functions  $f$ 's (representing posteriori knowledge) we want to discover
    - ①  $f$  is assumed to be parametrized by  $\mathbf{w}$
  - ② Define a **cost function**  $C(\mathbf{w}; \mathbb{X})$  (or functional  $C[f; \mathbb{X}]$ ) that measures “how good a particular  $f$  can explain the training data”

# General Machine Learning Steps

- ① Data collection, preprocessing (e.g., integration, cleaning, etc.), and exploration
  - ① Split a dataset into the training and testing datasets
- ② Model development
  - ① Assume a **model**  $\{f(\cdot; \mathbf{w})\}$  that is a collection of candidate functions  $f$ 's (representing posteriori knowledge) we want to discover
    - ①  $f$  is assumed to be parametrized by  $\mathbf{w}$
  - ② Define a **cost function**  $C(\mathbf{w}; \mathbb{X})$  (or functional  $C[f; \mathbb{X}]$ ) that measures “how good a particular  $f$  can explain the training data”
- ③ **Training**: employ an algorithm that finds the best (or good enough) function  $f^*(\cdot; \mathbf{w}^*)$  in the model that minimizes the cost function

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} C(\mathbf{w}; \mathbb{X})$$

# General Machine Learning Steps

- ① Data collection, preprocessing (e.g., integration, cleaning, etc.), and exploration
  - ① Split a dataset into the training and testing datasets
- ② Model development
  - ① Assume a **model**  $\{f(\cdot; \mathbf{w})\}$  that is a collection of candidate functions  $f$ 's (representing posteriori knowledge) we want to discover
    - ①  $f$  is assumed to be parametrized by  $\mathbf{w}$
  - ② Define a **cost function**  $C(\mathbf{w}; \mathbb{X})$  (or functional  $C[f; \mathbb{X}]$ ) that measures “how good a particular  $f$  can explain the training data”
- ③ **Training**: employ an algorithm that finds the best (or good enough) function  $f^*(\cdot; \mathbf{w}^*)$  in the model that minimizes the cost function

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} C(\mathbf{w}; \mathbb{X})$$

- ④ **Testing**: evaluate the performance of the learned  $f^*$  using the testing dataset

# General Machine Learning Steps

- ① Data collection, preprocessing (e.g., integration, cleaning, etc.), and exploration
  - ① Split a dataset into the training and testing datasets
- ② Model development
  - ① Assume a **model**  $\{f(\cdot; \mathbf{w})\}$  that is a collection of candidate functions  $f$ 's (representing posteriori knowledge) we want to discover
    - ①  $f$  is assumed to be parametrized by  $\mathbf{w}$
  - ② Define a **cost function**  $C(\mathbf{w}; \mathbb{X})$  (or functional  $C[f; \mathbb{X}]$ ) that measures “how good a particular  $f$  can explain the training data”
- ③ **Training**: employ an algorithm that finds the best (or good enough) function  $f^*(\cdot; \mathbf{w}^*)$  in the model that minimizes the cost function

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} C(\mathbf{w}; \mathbb{X})$$

- ④ **Testing**: evaluate the performance of the learned  $f^*$  using the testing dataset
- ⑤ Apply the model in the real world

# Example for Spam Detection

① Random split of your past emails and labels

① Training dataset:  $\mathbb{X} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_i$

② Testing dataset:  $\mathbb{X}' = \{(\mathbf{x}'^{(i)}, y'^{(i)})\}_i$

# Example for Spam Detection

## ① Random split of your past emails and labels

① Training dataset:  $\mathbb{X} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_i$

② Testing dataset:  $\mathbb{X}' = \{(\mathbf{x}'^{(i)}, y'^{(i)})\}_i$

## ② Model development

① **Model:**  $\{f : f(\mathbf{x}; \mathbf{w}) = \mathbf{w}^\top \mathbf{x}\}$

# Example for Spam Detection

## ① Random split of your past emails and labels

① Training dataset:  $\mathbb{X} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_i$

② Testing dataset:  $\mathbb{X}' = \{(\mathbf{x}'^{(i)}, y'^{(i)})\}_i$

## ② Model development

① **Model:**  $\{f : f(\mathbf{x}; \mathbf{w}) = \mathbf{w}^\top \mathbf{x}\}$

② **Cost function:**  $C(\mathbf{w}; \mathbb{X}) = \sum_i 1(\mathbf{w}; f(\mathbf{x}^{(i)}; \mathbf{w}) \neq y^{(i)})$



# Example for Spam Detection

- ① Random split of your past emails and labels
  - ① Training dataset:  $\mathbb{X} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_i$
  - ② Testing dataset:  $\mathbb{X}' = \{(\mathbf{x}'^{(i)}, y'^{(i)})\}_i$
- ② Model development
  - ① **Model:**  $\{f : f(\mathbf{x}; \mathbf{w}) = \mathbf{w}^\top \mathbf{x}\}$
  - ② **Cost function:**  $C(\mathbf{w}; \mathbb{X}) = \sum_i 1(\mathbf{w}; f(\mathbf{x}^{(i)}; \mathbf{w}) \neq y^{(i)})$
- ③ **Training:** to solve  $\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_i 1(\mathbf{w}; f(\mathbf{x}^{(i)}; \mathbf{w}) \neq y^{(i)})$

# Example for Spam Detection

- ① Random split of your past emails and labels
  - ① Training dataset:  $\mathbb{X} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_i$
  - ② Testing dataset:  $\mathbb{X}' = \{(\mathbf{x}'^{(i)}, y'^{(i)})\}_i$
- ② Model development
  - ① **Model:**  $\{f : f(\mathbf{x}; \mathbf{w}) = \mathbf{w}^\top \mathbf{x}\}$
  - ② **Cost function:**  $C(\mathbf{w}; \mathbb{X}) = \sum_i 1(\mathbf{w}; f(\mathbf{x}^{(i)}; \mathbf{w}) \neq y^{(i)})$
- ③ **Training:** to solve  $\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_i 1(\mathbf{w}; f(\mathbf{x}^{(i)}; \mathbf{w}) \neq y^{(i)})$
- ④ **Testing:** accuracy  $\frac{1}{|\mathbb{X}'|} \sum_i 1(\mathbf{x}'^{(i)}, y'^{(i)}; f(\mathbf{x}'^{(i)}; \mathbf{w}^*) = y'^{(i)})$

# Example for Spam Detection

- ① Random split of your past emails and labels
  - ① Training dataset:  $\mathbb{X} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_i$
  - ② Testing dataset:  $\mathbb{X}' = \{(\mathbf{x}'^{(i)}, y'^{(i)})\}_i$
- ② Model development
  - ① **Model:**  $\{f : f(\mathbf{x}; \mathbf{w}) = \mathbf{w}^\top \mathbf{x}\}$
  - ② **Cost function:**  $C(\mathbf{w}; \mathbb{X}) = \sum_i 1(\mathbf{w}; f(\mathbf{x}^{(i)}; \mathbf{w}) \neq y^{(i)})$
- ③ **Training:** to solve  $\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_i 1(\mathbf{w}; f(\mathbf{x}^{(i)}; \mathbf{w}) \neq y^{(i)})$
- ④ **Testing:** accuracy  $\frac{1}{|\mathbb{X}'|} \sum_i 1(\mathbf{x}'^{(i)}, y'^{(i)}; f(\mathbf{x}'^{(i)}; \mathbf{w}^*) = y'^{(i)})$
- ⑤ Use  $f^*$  to predict the labels of your future emails

# Example for Spam Detection

- ① Random split of your past emails and labels
  - ① Training dataset:  $\mathbb{X} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_i$
  - ② Testing dataset:  $\mathbb{X}' = \{(\mathbf{x}'^{(i)}, y'^{(i)})\}_i$
- ② Model development
  - ① **Model:**  $\{f : f(\mathbf{x}; \mathbf{w}) = \mathbf{w}^\top \mathbf{x}\}$
  - ② **Cost function:**  $C(\mathbf{w}; \mathbb{X}) = \sum_i 1(\mathbf{w}; f(\mathbf{x}^{(i)}; \mathbf{w}) \neq y^{(i)})$
- ③ **Training:** to solve  $\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_i 1(\mathbf{w}; f(\mathbf{x}^{(i)}; \mathbf{w}) \neq y^{(i)})$
- ④ **Testing:** accuracy  $\frac{1}{|\mathbb{X}'|} \sum_i 1(\mathbf{x}'^{(i)}, y'^{(i)}; f(\mathbf{x}'^{(i)}; \mathbf{w}^*) = y'^{(i)})$
- ⑤ Use  $f^*$  to predict the labels of your future emails
  - See Notation

# Outline

① What's Machine Learning?

② What's Deep Learning?


③ About this Course...

④ FAQ

# Deep Learning

- ML where an  $f(\cdot; \mathbf{w})$  has many (deep) layers

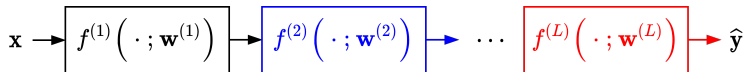
$$\hat{\mathbf{y}} = \color{red}{f^{(L)}}(\dots \color{blue}{f^{(2)}}(\color{black}{f^{(1)}}(\mathbf{x}; \mathbf{w}^{(1)}); \color{blue}{\mathbf{w}^{(2)}}) \dots; \color{red}{\mathbf{w}^{(L)}})$$



# Deep Learning

- ML where an  $f(\cdot; \mathbf{w})$  has many (deep) layers

$$\hat{\mathbf{y}} = f^{(L)}(\dots f^{(2)}(f^{(1)}(\mathbf{x}; \mathbf{w}^{(1)}); \mathbf{w}^{(2)}) \dots; \mathbf{w}^{(L)})$$

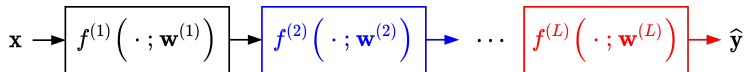


- Pros:
  - Learns to pre-process data automatically
  - Learns a complex function (e.g., visual objects to labels)

# Deep Learning

- ML where an  $f(\cdot; \mathbf{w})$  has many (deep) layers

$$\hat{\mathbf{y}} = f^{(L)}(\dots f^{(2)}(f^{(1)}(\mathbf{x}; \mathbf{w}^{(1)}); \mathbf{w}^{(2)}) \dots; \mathbf{w}^{(L)})$$



- Pros:
  - Learns to pre-process data automatically
  - Learns a complex function (e.g., visual objects to labels)
- Cons:
  - Usually needs large data to train a model well
  - Higher computation costs (for both training and testing)



# Outline

① What's Machine Learning?

② What's Deep Learning?

③ **About this Course...**

④ FAQ

# Target Audience

- *Senior undergraduate* and *graduate* CS students
  - Easy-to-moderate level of theory
  - Coding and engineering (in Python)
  - Clean datasets (small & large)

# Target Audience

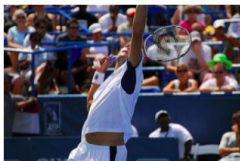
- *Senior undergraduate* and *graduate* CS students
  - Easy-to-moderate level of theory
  - Coding and engineering (in Python)
  - Clean datasets (small & large)
- No prior knowledge about ML is needed

# Topics Covered

- Supervised, unsupervised learning, and reinforcement learning

# Topics Covered

- Supervised, unsupervised learning, and reinforcement learning
- with *structural* output:



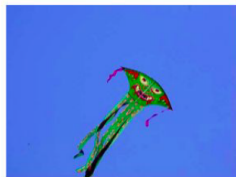
A man holding a tennis racquet on a tennis court.



Two pizzas sitting on top of a stove top oven



A group of young people playing a game of Frisbee



A man flying through the air while riding a snowboard

# Syllabus (Tentative)

- Part 1: math review (2 weeks)
  - Linear algebra
  - Probability & information theory
  - Numerical optimization

# Syllabus (Tentative)

- Part 1: math review (2 weeks)
  - Linear algebra
  - Probability & information theory
  - Numerical optimization
- Part 2: machine learning basics (3 weeks)
  - Learning theory
  - Parametric/non-parametric models
  - Experiment design

# Syllabus (Tentative)

- Part 1: math review (2 weeks)
  - Linear algebra
  - Probability & information theory
  - Numerical optimization
- Part 2: machine learning basics (3 weeks)
  - Learning theory
  - Parametric/non-parametric models
  - Experiment design
- Part 3: deep supervised learning (6 weeks)
  - Neural Networks (NNs), CNNs, RNNs



# Syllabus (Tentative)

- Part 1: math review (2 weeks)
  - Linear algebra
  - Probability & information theory
  - Numerical optimization
- Part 2: machine learning basics (3 weeks)
  - Learning theory
  - Parametric/non-parametric models
  - Experiment design
- Part 3: deep supervised learning (6 weeks)
  - Neural Networks (NNs), CNNs, RNNs
- Part 4: unsupervised learning (2 weeks)
  - Autoencoders, manifold learning, GANs

# Syllabus (Tentative)

- Part 1: math review (2 weeks)
  - Linear algebra
  - Probability & information theory
  - Numerical optimization
- Part 2: machine learning basics (3 weeks)
  - Learning theory
  - Parametric/non-parametric models
  - Experiment design
- Part 3: deep supervised learning (6 weeks)
  - Neural Networks (NNs), CNNs, RNNs
- Part 4: unsupervised learning (2 weeks)
  - Autoencoders, manifold learning, GANs
- Part 5: reinforcement learning (3 weeks)
  - Value/gradients policies, action/critics, reinforce RNNs

# Grading (Tentative)

- Prerequisite quiz: **15%**
  - ***On next Thu (9/21)***
  - ***You have to pass to be able to take this course: >70 or within top-70***
- Contests (× 4): **40%**
  - At the end of each part
- Assignments: **20%**
  - Come with the labs
- Final exam: **25%**
- Bonus: **6%**
  - Math labs (× 4)
  - Optional ML topics (× 2)

# Classes Info

- Lectures on Tue (2 hours)
  - Concepts & theories
  - with companion videos
- Labs on Thu (1 hour)
  - Implementation (in Python) & engineering topics
- TA time: 4:20pm–5:30pm on Thu at Delta 729
- More info can be found in the [course website](#)

# Outline

① What's Machine Learning?

② What's Deep Learning?

③ About this Course...

**④ FAQ**

# FAQ (1/2)

*Q: Should we team up for the contests?*

**A:** Yes, 2~4 students per team

# FAQ (1/2)

*Q: Should we team up for the contests?*

*A:* Yes, 2~4 students per team

*Q: Which GPU card should I buy?*

*A:* Nvidia GTX 1060 or above; 1050 Ti (4G RAM) minimal

# FAQ (1/2)

*Q: Should we team up for the contests?*

*A:* Yes, 2~4 students per team

*Q: Which GPU card should I buy?*

*A:* Nvidia GTX 1060 or above; 1050 Ti (4G RAM) minimal

*Q: Do we need to attend the classes?*

*A:* No, as long as you can pass. But you have attendance bonus...



# FAQ (1/2)

*Q: Should we team up for the contests?*

*A: Yes, 2~4 students per team*

*Q: Which GPU card should I buy?*

*A: Nvidia GTX 1060 or above; 1050 Ti (4G RAM) minimal*

*Q: Do we need to attend the classes?*

*A: No, as long as you can pass. But you have attendance bonus...*

*Q: Is this a light-loading course or heavy-loading one?*

*A: Should be **very heavy** to most students. Please **reserve your time***

# FAQ (2/2)

*Q: What's the textbook?*

*A:* No formal textbook. But if you need one, read the [Deep Learning](#) book

# FAQ (2/2)

*Q: What's the textbook?*

*A:* No formal textbook. But if you need one, read the [Deep Learning](#) book

*Q: Why some sections are marked with "\*" or "\*\*" in the slides?*

*A:* The mark "\*" means "can be skipped for the first time reader," and  
" \*\*" means "materials for reference only"

# TODO

- Assigned reading:
  - Calculus
  - Get your feet wet with Python

# Reference I

- [1] D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. Van Den Driessche, J. Schrittwieser, I. Antonoglou, V. Panneershelvam, M. Lanctot, et al.  
Mastering the game of go with deep neural networks and tree search.  
*Nature*, 529(7587):484–489, 2016.